

# LDA

Aidan Boland

2015-10-15

## Latent Dirichlet Allocation

This document is a summary of Latent Dirichlet Allocation, a method first presented in [Blei et al., 2002](#).

### Notation

The data is composed as follows,

- A vocabulary of  $V$  possible words  $(w^1, \dots, w^V)$ .
- A corpus of  $M$  individual documents,  $\mathbf{D} = (\mathbf{w}_1, \mathbf{w}_2, \dots, \mathbf{w}_M)$ .
- Document  $d$  is composed of  $N_d$  words,  $\mathbf{w}_d = (w_{d,1}, w_{d,2}, \dots, w_{d,N_d})$ .
- There are  $k$  topics  $(1, \dots, k)$ .

For consistency the following letters will be used as sub and super scripts,

- $i$  for topic  $(1, \dots, k)$ .
- $j$  for word in vocabulary  $(1, \dots, V)$ .
- $n$  for word in document  $(1, \dots, N_d)$ .
- $d$  for document  $(1, \dots, M)$ .

### Generating a document

The following steps show how document  $\mathbf{w}_d$  is generated.

1. Choose the number of words in the document,  $N_d \sim Po(\xi)$ .
2. Choose the probabilities of the topics for the document,  $\theta_d \sim Dir(\alpha)$ . ( $\theta$  is a  $1 \times k$  vector)
3. For each of the  $N_d$  words  $w_{d,n}$ ,
  - Choose a topic  $z_{d,n} \sim Mult(\theta_d)$ , ( $z_{d,n}$  is a single value from 1 to  $k$ , the topic of word  $n$  in doc  $d$ )
  - Choose a word  $w_{d,n} \sim Mult(\phi_{z_{d,n}})$ ,

where  $\phi \sim Dir(\beta)$ . The choices of the hyperparameters  $\alpha$  and  $\beta$  will be discussed further in the next section. Figure 3 shows a graphical representation of LDA.

### Parameters

The word probabilities are parameterised by the  $k \times V$  matrix  $\phi$ ,

$$P(w_{d,n} = j | z_{d,n} = i) = \phi_{j,i}$$

The topic probabilities for each document are parameterised by the vector  $\theta_d$

$$P(z_{d,n} = j) = \theta_{d,j}$$




Figure 1: Graphical model for LDA.

Latent dirichlet as presented in [Blei et al., 2002](#) is a problem of maximising the following,

$$P(D|\phi, \alpha) = \int P(D|\phi, \theta)P(\theta|\alpha)d\theta,$$

this can be achieved using variational inference and the EM algorithm. (note, in the Blei paper they do not put a prior on  $\phi$ )

### Gibbs sampling

An alternative efficient procedure is [collapsed Gibbs sampling \(alternative paper\)](#) where  $\theta$  and  $\phi$  are marginalised out and only the latent variables  $\mathbf{z}$  are sampled. After the sampler has burned-in we can calculate an estimate of  $\theta$  and  $\phi$  given  $\mathbf{z}$ .

The posterior distribution of interest in this case is,

$$P(\mathbf{z}|D) = \frac{P(D, \mathbf{z})}{\sum_{\mathbf{z}} P(D, \mathbf{z})} \propto P(D|\mathbf{z})P(\mathbf{z})$$

The joint distribution,  $P(D, \mathbf{z})$ , can be found through integrating

$$P(\theta, \phi, \mathbf{z}, D|\alpha, \beta) = P(D|\mathbf{z}, \phi)P(\phi|\beta)P(\mathbf{z}|\theta)P(\theta|\alpha)$$

with respect to  $\phi$  and  $\theta$ .

Integrating with respect to  $\phi$  gives the following,

$$p(D|\mathbf{z}) = \prod_{i=1}^K \left[ \frac{\Gamma\left(\sum_{j=1}^V \beta_{i,j}\right)}{\prod_{j=1}^V \Gamma(\beta_{i,j})} \times \frac{\prod_{j=1}^V \Gamma(C_{j,\cdot}^i + \beta_{i,j})}{\Gamma\left(\sum_{j=1}^V C_{j,\cdot}^i + \beta_{i,j}\right)} \right] \quad (1)$$

while integrating with respect to  $\theta$  gives,

$$p(\mathbf{z}) = \prod_{d=1}^M \left[ \frac{\Gamma\left(\sum_{j=1}^K \alpha_j\right)}{\prod_{j=1}^K \Gamma(\alpha_j)} \times \frac{\prod_{j=1}^K \Gamma(C_{j,d} + \alpha_j)}{\Gamma\left(\sum_{j=1}^K (C_{j,d} + \alpha_j)\right)} \right] \quad (2)$$

Where  $C_{j,d}^{i,-(d,n)}$  is the count of words  $w^j$  that have topic  $i$  in document  $\mathbf{w}_d$  not including the word  $w_{d,n}$ . The notation  $\cdot$  is a sum over that subscript, e.g.  $C_{j,\cdot}^{i,-(d,n)}$  is the number of words  $w^j$  that are assigned to topic  $j$  in the full corpus not including the word  $w_{d,n}$ .

The full conditional distribution  $P(z_{d,n} = i | \mathbf{z}_{-(d,n)}, w_{d,n} = j, D, \alpha, \beta)$  can be found using equations (1) and (2),

$$P(z_{d,n} = i | \mathbf{z}_{-(d,n)}, w_{d,n} = j, D, \alpha, \beta) \propto \frac{\beta_{i,j} + C_{j,\cdot}^{i,-(d,n)}}{\beta_{i,\cdot} + C_{\cdot,\cdot}^{i,-(d,n)}} \frac{\alpha_i + C_{\cdot,d}^{i,-(d,n)}}{\alpha_{\cdot} + C_{\cdot,d}^{i,-(d,n)}}$$

The labels in the corpus can be updated with Gibbs sampling using full conditional distribution.

## Implementation

The following code can be used to update the labels using the collapsed Gibbs sampler as described above. The first function creates documents given a set of priors.

```
library(MCMCpack) # for dirichlet dist, alternative is library(gtools)

GenerateLDA <- function(alpha, beta, xi, Vocab, Ndoc){
  # function to generate a single document
  # arguments:
  #   alpha, k x 1 concentration paramter for topics, (take input as single value for now)
  #   beta, k x V each row represents the concentration parameter for words in topic k
  #   xi, parameter to simulate length of documents
  #   Vocab, vector of size V containing all possible words
  #   ndoc, number of documents to simulate
  #
  # todo:
  #   change output to document term matrix rather than list....

  alpha <- rep(alpha, nrow(beta))

  if(dim(beta)[2] != length(Vocab))
    stop("Incorrect dimension for Beta!!")

  docs <- list() # blank list to store words, documents diff size so array
                # is not suitable, could use DTM alternatively...
  z_out <- NULL
  for(i in 1:Ndoc){
    N <- rpois(1, lambda = xi) # Choose the length of the ith doc
    theta <- rdirichlet(n = 1, alpha = alpha)

    z <- rmultinom(1, 1, prob = theta) # Choose a topic for 1st word in doc i
    z_out <- c(z_out, which(z == 1))
    phi <- rdirichlet(n = 1, alpha = beta[which(z == 1),]) # generate probs
    docs[[i]] <- Vocab[which(rmultinom(1, 1, prob = phi) == 1)] # Initialise ith doc

    for(j in 2:N){
      z <- rmultinom(1, 1, prob = theta) # Choose a topic for jth word in doc i
      z_out <- c(z_out, which(z == 1))
      phi <- rdirichlet(n = 1, alpha = beta[which(z == 1),])
      docs[[i]] <- c(docs[[i]], Vocab[which(rmultinom(1, 1, prob = phi) == 1)])
    }
  }
  list(docs = docs, z = z_out)
}
```

```

my_topic_counts <- function(dtm, z, n_topics){
  # Function to calculate topic counts for words and documents
  # arguments:
  #   dtm is a document term matrix in triplet form
  #   dtm$i denotes the document
  #   dtm$j denotes the word
  #   dtm$v is the number of occurrences of word j in doc i
  #   z is a set of topic labels for each word

  word <- rep(dtm$j, dtm$v)
  doc <- rep(dtm$i, dtm$v)

  word_count <- array(0, c(n_topics, ncol(dtm)))
  doc_count <- array(0, c(n_topics, nrow(dtm)))

  for(w in 1:length(doc)){
    word_count[z[w], word[w]] <- word_count[z[w], word[w]] + 1
    doc_count[z[w], doc[w]] <- doc_count[z[w], doc[w]] + 1
  }
  list(word_count, doc_count)
}

my_lda_gibbs <- function(dtm, n_topic = 2, iterations = 10){

  n_vocab <- ncol(dtm) # Number of unique words
  alpha <- rep(1, n_topic) # prior on theta
  beta <- array(1, c(n_topic, n_vocab)) # prior on phi
  theta <- array(0, c(nrow(dtm), n_topic))
  phi <- array(0, c(n_topic, n_vocab))

  doc <- rep(dtm$i, dtm$v)
  word <- rep(dtm$j, dtm$v)

  z <- array(0, c(iterations + 1, length(doc)))
  z[1, ] <- sample(1:n_topic, length(doc), replace = TRUE) # choose initial labels
  if(length(unique(z[1, ])) != n_topic)
    stop("Number of topics in z vector do not match argument!\n")

  initial_counts <- my_topic_counts(dtm, z[1, ], n_topics = n_topic)
  word_count <- initial_counts[[1]]
  doc_count <- initial_counts[[2]]
  topic_prob <- array(0, n_topic)

  for(iter in 1:iterations){
    for(w in 1:length(word)){
      # Loop over each word
      for(j in 1:n_topic){
        if(z[iter, w] == j){
          topic_prob[j] <- prod(sum(beta[j, word[w]], word_count[j, word[w]], -1),
                                sum(alpha[j, doc[w]], doc_count[j, doc[w]], -1)) /
            prod(sum(beta[j, ], word_count[j, ], -1),
                  (sum(alpha, doc_count[, doc[w]], -1)))
        }else{

```

```

        topic_prob[j] <- prod(sum(beta[j, word[w]], word_count[j, word[w]]),
                             sum(alpha[j], doc_count[j, doc[w]])) /
        prod(sum(beta[j, ], word_count[j, ]),
              sum(alpha, doc_count[, doc[w]]))
    }
}
z[iter + 1, w] <- sample.int (n_topic, size = 1, prob = topic_prob)

if(z[iter + 1, w] != z[iter, w]){
  word_count[z[iter, w], word[w]] <- word_count[z[iter, w], word[w]] - 1
  doc_count[z[iter, w], doc[w]] <- doc_count[z[iter, w], doc[w]] - 1
  word_count[z[iter + 1, w], word[w]] <- word_count[z[iter + 1, w], word[w]] + 1
  doc_count[z[iter + 1, w], doc[w]] <- doc_count[z[iter + 1, w], doc[w]] + 1
}
}
# Old code when counts updated after each full sweep of words
#update_counts <- my_topic_counts(dtm, z[iter + 1, ], n_topic)
#word_count <- update_counts[[1]]
#doc_count <- update_counts[[2]]
}
for(j in 1:n_topic)
  phi[j, ] <- (beta[j, ] + word_count[j, ]) / (sum(beta[j, ]) + sum(word_count[j, ]))
for(d in 1:nrow(dtm))
  theta[d, ] <- (alpha + doc_count[, d]) / (sum(alpha) + sum(doc_count[, d]))

list(z, t(phi), theta)
}

```

## Testing the functions

We can now implement these functions using a sample set of words. The corpus was created using a specific set of  $\beta$  prior values.

```

my_vocab <- c("car", "engine", "exhaust", "wheel",
             "milk", "cream", "dairy", "yogurt",
             "coke", "water", "juice", "coffee", "drink", "bottle", "can")
n_topics <- 3
# my_beta <- matrix(1, nrow = n_topics, ncol = length(my_vocab))
my_beta <- matrix(c(10, 10, 10, 10, 0.1, 0.1, 0.1, 0.1, 0.1, 0.1, 0.1, 0.1, 0.1, 0.1, 0.1,
                  0.1, 0.1, 0.1, 0.1, 10, 10, 10, 10, 0.1, 0.1, 0.1, 0.1, 0.1, 0.1, 0.1,
                  0.1, 0.1, 0.1, 0.1, 0.1, 0.1, 0.1, 0.1, 10, 10, 10, 10, 10, 10),
                  byrow = T, nrow = n_topics, ncol = length(my_vocab))
my_corpus <- GenerateLDA(alpha = 1, beta = my_beta, xi = 10, Vocab = my_vocab, Ndoc = 100)
library(tm) # Library to create DTM
my_dtm <- DocumentTermMatrix(Corpus(VectorSource(my_corpus$docs)))

```

The corpus was then analysed using a pre-built function from the *topicmodels* package (LDA) and the function created above (`my_lda_gibbs`). The priors for the Gibbs steps were left vague.

```

library(topicmodels)
time_pacakge <- system.time(
  control_gibbs <- LDA(my_dtm, k = 3, method="Gibbs", iterations = 3000, burn.in=1000)

```

```
)

time_mine <- system.time(
  my_gibbs <- my_lda_gibbs(my_dtm, n_topic = 3, iterations = 3000)
)
```

The LDA function took 0.11 seconds to run while the `my_lda_gibbs` function took 89.51 seconds. The LDA function is optimised to run in `c`, this is where it gains its time advantage. One problem is that neither of the functions seem to converge well. The tables below compare the 2 algorithms with the true values for each word in the corpus.

```
my_gibbs_z <- apply(my_gibbs[[1]][-1000, ], 2, median)
table(z = my_corpus$z, my_gibbs = my_gibbs_z)
```

	my_gibbs		
z	1	2	3
1	65	74	161
2	75	204	63
3	191	83	67

```
table(z = my_corpus$z, control_gibbs@z)
```

z	1	2	3
1	170	68	62
2	68	195	79
3	83	81	177

```
table(package_gibbs = control_gibbs@z, my_gibbs = my_gibbs_z)
```

	my_gibbs		
package_gibbs	1	2	3
1	29	4	288
2	1	342	1
3	301	15	2

Seems to do ok, not fantastic. There may be a label matching problem but this not too big of a deal, it is a standard problem in clustering. What is nice is that both algorithms seem to get the same results (last table). This indicates that the functions created in this document are correct.

## Seeded topics

One simple solution to incorporating seed words into the LDA model is presented in [Jagarlamundi et. al](#) [?](#). A set of seed words can be provided by the user, these are used to help the model learn about the topics. Jagarlamundi et. al build a model which uses the seed words to improve both the topic-word and document-topic probability distributions. Following the paper both of these models are presented separately first (Model 1 and Model 2) and then combined.

## Word-Topic Distributions (Model 1)

This model chooses words from two Multinomial distributions: a **seed topic' distribution** and a **regular topic' distribution**. The seed topic distribution is constrained to only generate words from a corresponding seed set. The regular topic distribution may generate any word including seed words.

1. Choose regular topic  $\phi_k^r \sim Dir(\beta_r)$
2. Choose seed topic  $\phi_k^s \sim Dir(\beta_s)$
3. Choose  $\pi_k \sim Beta(1, 1)$

Document  $\mathbf{w}_d$  is then generated as follows.

1. Choose the number of words in the document,  $N_d \sim Po(\xi)$ .
2. Choose the probabilities of the topics for the document,  $\theta_d \sim Dir(\alpha)$ . ( $\theta$  is a  $1 \times k$  vector)
3. For each of the  $N_d$  words  $w_{d,n}$ ,
  - Choose a topic  $z_{d,n} \sim Mult(\theta_d)$ , ( $z_{d,n}$  is a single value from 1 to  $k$ , the topic of word  $n$  in doc  $d$ )
  - Select an indicator  $x_{d,n} \sim Bern(\pi_{z_{d,n}})$
  - If  $x_{d,n} = 0$  choose a word  $w_{d,n} \sim Mult(\phi_{z_{d,n}}^s)$ ,
  - If  $x_{d,n} = 1$  choose a word  $w_{d,n} \sim Mult(\phi_{z_{d,n}}^r)$ ,

where  $\phi \sim Dir(\beta)$ . The choices of the hyperparameters  $\alpha$  and  $\beta$  will be discussed further in the next section. Figure 3 shows a graphical representation of LDA.




Figure 2: Graphical model for Model 1.

## Document-Topic Distributions (Model 2)

Very high level idea is that documents which contain a word from one of the seed topics have a higher probability of coming from this topic. ....sounds a bit simple

1. For each topic choose  $\phi_j \sim Dir(\beta)$
2. For each seed topic choose a group-topic distribution  $\psi_j \sim Dir(\alpha)$

Document  $\mathbf{w}_d$  is then generated as follows.

1. Choose binary vector  $b$  (vector of length  $K$  indicating whether seed topic  $j$  is in document)
  - when generating documents this can be chosen randomly

2. Choose document-group distribution  $\xi_d \sim \text{Dir}(\tau b)$
3. Choose a group variable  $g \sim \text{Mult}(\xi_d)$
4. Choose  $\theta_d \sim \text{Dir}(\psi_g)$
5. For each of the  $N_d$  words  $w_{d,n}$ ,
  - Choose a topic  $z_{d,n} \sim \text{Mult}(\theta_d)$ ,
  - Choose a word  $w_{d,n} \sim \text{Mult}(\phi_{z_{d,n}})$ .

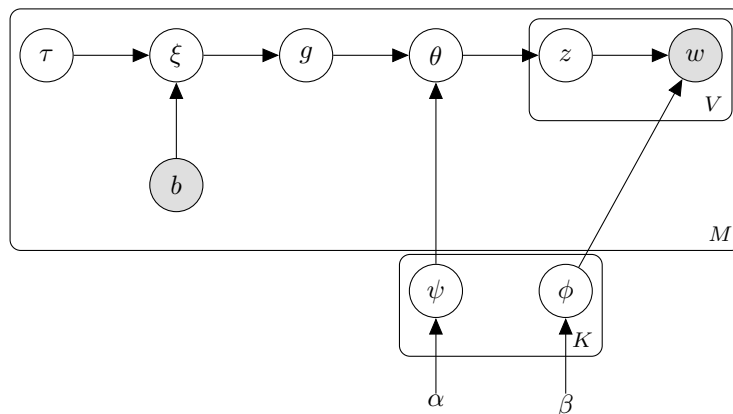


Figure 3: Graphical model for Model 2.

## Hierarchical LDA

Hierarchically Supervised Latent Dirichlet Allocation [Perotte et. al 11.](#)